

Faster Python Programs - Measure, don't Guess

A Tutorial at PyCon 2019

May 2, 2019

Cleveland, OH, USA

author: Dr.-Ing. Mike Müller

email: mmueller@python-academy.de

twitter: @pyacademy

version: 4.1

About Me

- Enthusiastic Python user since 1999
- Python trainer since 2004
- CEO of Python Academy
- Chair of EuroSciPy 2008, 2009, PyCon DE 2011, 2012 and EuroPython 2014
- Still involved in organizing Python conferences
- Chair of PySV e.V. (German equivalent to the PSF)

About You

- How many years of Python?
- Python 2 or 3?
- Do you do number crunching?
- Are you problems CPU bound?
- Or memory bound?
- Or IO bound?

Overview

- Hands-on sessions
- IPython Notebook recommended (standard Python and editor works too)
- PDF handout
- Download of examples
- Exercises after each topics
- 20 minutes break
- Please give feedback via Guidebook or Web (<https://www.surveymonkey.com/...>)

How Fast is Fast Enough?

- Python is very high-level
- Slow for some types of computations
- This may or may not be a problem
- Before starting to optimize consider the costs
- Optimized code may need more effort to develop and maintain
- Balance between speed of development and speed of program execution

Optimization Guidelines

Premature optimization is the root of all evil.

D. Knuth or C. A. R. Hoare or folklore (attributions seems not totally clear)

Correct First - Faster Later

- Make sure your program works correctly
- Before you start thinking about optimization
- Never optimize before the program produces the desired results

The Prize of Optimization

- Optimization often comes with a price
- It tends to make your code less readable
- More maintenance effort
- Readability and maintainability can suffer
- Python's excellent readability is not for free in terms of runtime
- You cannot always solve this contradiction between both goals

General Guidelines

- ... for optimization

1. Make sure your program is really too slow

- Do you really need more performance?
- Are there any other slowdown factors such as network traffic or user input that have more impact on speed?
- Does it hurt if the program runs slowly?

2. Don't optimize as you go

- Don't waste time before you are certain that you will need the additional speed.

3. Only realistic use cases and user experience should be considered

- Maximal expected number of users
- How often the program is used

4. Search for external causes for slowdown

- Network connections
- Database access
- Calls to system functions

5. Architecture can be essential for performance

- Is it appropriate?

6. Are there any bugs that slow down the program?

- Don't optimize with bugs
- The bug might be the bottleneck

7. If the program is really too slow, find the bottlenecks by profiling

- That is what this tutorial is about

8. Always check the result of optimization

- Use unit tests
- With high coverage

9. Go for algorithms first

- Typically most performance gains come from better algorithms
- Find the big-O notation of your algorithm

10. Throw better hardware on it

- Often hardware is cheaper than programmer time
- You often easily add memory (avoid swapping)
- Use the fastest processor you can get